

▽ giza: A Related Set of Bash Expansions of Relevance to AI

Arash Abadpour - arash@abadpour.com

November 17, 2023

Abstract

wip

In this paper, we discuss the Mathematics of building Machine Vision AI systems in Linux. We review the general challenge of translating the description of an AI operation in human language into a human-readable, machine-executable script. We select multiple Machine Vision AI challenges that we first describe in human language. Then, in each case, we build the language to convert the description in human language into one or more scripts we execute on machines. We use AWS SageMaker ¹ for development and training and AWS Batch ² for inference and discuss API calls. The main contribution of this paper is a mathematical framework for building an AI language for a practical use-case in Machine Vision. We hope that researchers in other fields of AI use and extend this framework in their disciplines. We present a reference implementation ³ of this framework and multiple use-cases ⁴ - *revision-1.80.1*

Contents

1	Introduction	3
2	Expansions	4
2.1	Command Substitution	4
2.1.1	@cache	4
2.1.2	@ref	4
2.1.3	@tag	4
2.1.4	@relations	5
2.2	Options	5
2.3	Objects	5
2.3.1	Object Metadata	5
2.3.2	@select and object references	5
2.3.3	@upload	6
2.3.4	@download	6
2.3.5	@list	6
2.4	--<keyword> <value>	6
2.5	Prefixing	6
2.6	@docker	6
2.7	<command> help	6
2.8	@init	6
2.9	The Core	6
2.10	Plugins	6
2.11	Scripts	6
2.12	@seed	6
2.13	@start	6
2.14	@conda	6

¹<https://aws.amazon.com/sagemaker/>

²<https://aws.amazon.com/batch/>

³<https://github.com/kamangir/awesome-bash-cli>, *awesome-bash-cli*, *abcli*.

⁴bird watching in downtown Vancouver with AI, <https://github.com/kamangir/Vancouver-Watching>, *Vancouver-Watching*, *vanwatch*.

2.15	@git	6
2.16	@<keyword>	6
3	Examples	7
3.1	hubble	7
3.2	OpenAI code generation	7
3.3	roofAI	7
3.3.1	roofAI dataset	7
3.3.2	roofAI semseg train	7
3.3.3	roofAI semseg predict	7
3.4	Vancouver-Watching (vanwatch)	8
3.4.1	vanwatch discover	8
3.4.2	vanwatch ingest	8
3.4.3	vanwatch process	8
A	Concepts	9
A.1	Machines and Shells	9
A.2	Operators	9
A.3	Commands	9
A.3.1	Callables	10
A.3.2	Command Templates	10
A.4	Computational Model	10
B	Background and Context	11

1 Introduction

Bash is a “Unix shell and command language first released in 1989 that has been used as the default login shell for most Linux distributions”⁵. A shell is a “macro processor that executes commands”⁶, where “macro processor means functionality where text and symbols are expanded to create larger expressions” (same reference). There are seven kinds of expansions in Bash⁷.

*Brace Expansion*⁸ is the first and the quickest to explain,

```
> bash$ echo a{d,c,b}e
ade ace abe
```

*Tilde Expansion*⁹ relates to words that begin with an unquoted tilde character (~). *Parameter and Variable Expansion*¹⁰ enable the use of variables, as `${variable}`, as well as more elaborate pattern matching forms such as `${parameter/#pattern/string}`. *Command Substitution* “allows the output of a command to replace the command itself”¹¹. *Arithmetic expansion*¹² enables arithmetic operations using the form `$((expression))` and *Word Splitting*¹³ governs the splitting of the command to words. Finally, *Filename Expansion*¹⁴ enables the familiar wildcard reference to filenames using ‘*’ and ‘?’. In Section 2 we propose a set of relevant expansions to AI operations that are implemented using Python¹⁵.

This work, first, proposes several novel Bash expansions based on command substitution 2.1. Then, using typed positional arguments, we propose *options* 2.2 and *objects* 2.3, including *object metadata* ??. Then, we discuss argument injection as a suffix 2.4 and *Prefixing* 2.5 to transform `<command>` to `<prefix> <command>`, as used by `@docker` 2.6, for example. The next two expansions, `<command> help` 2.7 and `@init <args>` 2.8 are suggestions to reduce the cognitive load of the user.

Then, we discuss the `<core>` 2.9, which is the callable that is `source`’d in a startup file¹⁶. The `plugins` 2.10 enable the addition of branches to the syntax, and `scripts` 2.11 implement the last mile. We then discuss the `@seed` 2.12; the notion that code generates code that is transferred into another machine through the clipboard, a key, or a `scp`¹⁷-style protocol to terraform the machine and start a process A.3. `@start` 2.13 is a necessity; the first intelligent command to start the day with. `@start` behaves according to the machine it runs on and other aspects of the state A.1. On a MacBook, `@start` logs in and starts an `ssh` session to the default machine. On that machine, `@start` starts the docker container. Note that the specific action of `@start` is decided by its immediate user¹⁸. We then discuss practical expansions for `@conda` 2.14, `@git` 2.15, and `wip`. Finally, we discuss the `@` 2.16.

⁵[https://en.wikipedia.org/wiki/Bash_\(Unix_shell\)](https://en.wikipedia.org/wiki/Bash_(Unix_shell))

⁶https://www.gnu.org/software/bash/manual/html_node/What-is-a-shell_003f.html

⁷https://www.gnu.org/software/bash/manual/html_node/Shell-Expansions.html

⁸https://www.gnu.org/software/bash/manual/html_node/Brace-Expansion.html

⁹https://www.gnu.org/software/bash/manual/html_node/Tilde-Expansion.html

¹⁰https://www.gnu.org/software/bash/manual/html_node/Shell-Parameter-Expansion.html

¹¹https://www.gnu.org/software/bash/manual/html_node/Command-Substitution.html

¹²https://www.gnu.org/software/bash/manual/html_node/Arithmetic-Expansion.html

¹³https://www.gnu.org/software/bash/manual/html_node/Word-Splitting.html

¹⁴https://www.gnu.org/software/bash/manual/html_node/Filename-Expansion.html

¹⁵<https://github.com/kamangir/awesome-bash-cli>

¹⁶https://www.gnu.org/software/bash/manual/html_node/Bash-Startup-Files.html

¹⁷<https://linux.die.net/man/1/scp>

¹⁸the tool is adapted to the tool user.

2 Expansions

2.1 Command Substitution

A symbolic description of command substitution is that if `<part-1>$(<sub-command>)<part-2>` is encountered, then `<sub-command>` is executed and its outcome, `<outcome>`, is used to generate the updated command as `<part-1><outcome><part-2>`, which is then executed. This process happens numerous and hierarchically ¹⁹.

Here is an example from `roofAI` 3.3,

```
roofAI semseg predict \  
  profile=FULL,upload \  
  $(@ref roofAI_semseg_model_AIRS_full_v2) \  
  $(@ref roofAI_ingest_AIRS_v2) \  
  $(@timestamp)
```

Here, `@ref <keyword>` reads the value of `<keyword>` from the `cache` 2.1.1 and `@timestamp` generates a unique timestamp for use as an `<object-name>`. Collectively, this command runs the “Pytorch Segmentation Model” ²⁰ that is cached as `roofAI_semseg_model_AIRS_full_v2` on the dataset that is cached as `roofAI_ingest_AIRS_v2` and uploads the results in a uniquely named object 2.3. Tags 2.1.3 and relations 2.1.4 are other object metadata 2.3.1 that are relevant to this expansion.

Command substitution is useful for generating the command components through Python or Bash. For example, in the AWS Open Data Registry ²¹ there is the notion of datasets, such as `hst` ²² for Hubble Space Telescope and metadata the dataset is maintained in `yaml` files a git repository ²³.

For example, here is the command to access `ibrma2f2q_drc.jpg` in object `public/ibrm/ibrma2f2q` in the dataset `hst`,

```
aws s3 cp $auth $s3_uri$filename $path
```

Here, `$auth` and `$s3_uri` are generated as,

```
local auth=$(abcli_hubble_get auth $dataset_name)  
local s3_uri=$(abcli_hubble_get s3_uri $dataset_name $hubble_object_name)
```

Here, `abcli_hubble_get` is a Bash wrapper around a Python call.

```
function abcli_hubble_get() {  
  python3 -m hubble get \  
    --what "$1" \  
    --dataset_name "$2" \  
    --object_name "$3" \  
    "${@:4}"  
}
```

These were examples of command substitution with Bash and Python calls from `hubble` 3.1 which sits on the `abcli` ²⁴ core 2.9.

2.1.1 @cache

2.1.2 @ref

2.1.3 @tag

An object can have many tags. A tag is a boolean or valued property of the object.

¹⁹https://www.gnu.org/software/bash/manual/html_node/Command-Substitution.html

²⁰https://github.com/qubvel/segmentation_models.pytorch

²¹<https://registry.opendata.aws/>

²²<https://registry.opendata.aws/hst/>

²³<https://github.com/aws-labs/open-data-registry/blob/main/datasets/hst.yaml>

²⁴<https://github.com/kamangir/awesome-bash-cli>

```
@tags get <object-name>
@tags get <object-name> <keyword>

@tags search <keyword=value,~that,this>

@tags set <object-name> <keyword=value,~that,this>
```

2.1.4 @relations

Two objects can be related in many ways.

```
@relations get <object-name-1> <object-name-2>
@relations get <object-name-1> <object-name-2> <relation>
@relations get <object-name-1>
@relations get <object-name-1> <relation>

@relations search <relation>

@relations set <object-name-1> <object-name-2> <relation>
```

2.2 Options

An *options* is a string representation of a dictionary, such as <keyword>=<value>,<keyword>,-<keyword>.

2.3 Objects

Objects are accessible on every machine by their *name*, and an object may be *selected* 2.3.2. Objects may be uploaded 2.3.3 and downloaded 2.3.4, fully or partially, or listed 2.3.5. Objects have metadata ??.

Commands A.3 consume and generate objects. If the object name is not provided or is given as “-”, then an object with a unique name is created and used. Otherwise, the object is either provided by name, as in <object-name>, or by reference, as in one of “.”, “..” or “...” 2.3.2.

In its most basic form, an object is an S3 bucket ²⁵.

2.3.1 Object Metadata

metadata is information about Objects 2.3, such as their *tags* 2.1.3 and relations 2.1.4.

2.3.2 @select and object references

```
@select object <object-name>
@select <type> <type-name>
```

When an object is selected, ‘.’ represents <object-name_i>. Similarly, ‘..’, ‘...’, and so on, as deep as needed, refer to the previously selected object and the one before that.

Commands default the objects they consume and modify to ‘.’, ‘..’, and so on. Because the commands in a script use the same objects, selecting the objects enables their names to be omitted in a script.

An object may have a ‘type’, for example, ‘model’ or ‘dataset’. Commands that consume objects specify a ‘type’ for the argument. This enables the user to simultaneously select different types of objects and run commands on them. Here is an example from *hubble* ²⁶, wherein the user selects an object and then a hubble dataset and an object in that dataset to download in the object.

```
abcli select; open .
hubble select dataset hst
hubble select object public/u4ge/u4ge0106r
hubble download ~dryrun
```

²⁵<https://aws.amazon.com/s3/>

²⁶<https://github.com/kamangir/hubble>

2.3.3 @upload

@upload object <object-name> [filename=<filename>]

2.3.4 @download

@download object <object-name> [filename=<filename>,open]

2.3.5 @list

2.4 --<keyword> <value>

2.5 Prefixing

2.6 @docker

2.7 <command> help

When the command ends with `help`, a help message prints.

2.8 @init

git pull

2.9 The Core

The *core* is a callable that loads the plugins and terraforms the machine.

2.10 Plugins

A plugin generally defines one or more callables.

2.11 Scripts

The script is a bash script ²⁷.

2.12 @seed

2.13 @start

2.14 @conda

2.15 @git

git pull

2.16 @<keyword>

²⁷https://en.wikipedia.org/wiki/Shell_script

3 Examples

3.1 hubble

`hubble` is a ∇ giza callable ?? to access and process Hubble Space Telescope imagery and other datasets on AWS Open Data Registry ²⁸.

3.2 OpenAI code generation

Experiments w/ the OpenAI API ²⁹. wip

3.3 roofAI

`roofAI` ³⁰ is a callable that,

1. terraforms the machine and the shell [2.14](#).
2. ingests and reviews datasets [3.3.1](#).
3. train semantic segmentation ³¹ models [3.3.2](#).
4. runs semantic segmentation predictions [3.3.3](#).

3.3.1 roofAI dataset

`roofAI dataset ingest source=<source> <object-name>` ingests a dataset from `source` into `<object-name>` and tags it for future discovery. `roofAI dataset review - <object-name>` reviews the dataset in `<object-name>`.

```
roofAI dataset ingest \  
  [source=AIRS|CamVid,register,suffix=<v1>] \  
  <object-name> \  
  [<args>]
```

```
roofAI dataset review \  
  [-] \  
  <object-name> \  
  [<args>]
```

3.3.2 roofAI semseg train

`semseg train - <dataset-object-name> <model-object-name>` trains a model on `<dataset-object-name>` that it stores in `<model-object-name>` and tags for future discovery.

```
semseg train \  
  [device=cpu|cuda,register,suffix=<v1>] \  
  <dataset-object-name> \  
  <model-object-name> \  
  [<args>]
```

3.3.3 roofAI semseg predict

`semseg predict - <model-object-name> <dataset-object-name>` runs a prediction on the dataset `<dataset-object-name>` using the model `<model-object-name>`.

²⁸<https://github.com/kamangir/hubble>

²⁹<https://github.com/kamangir/openai>

³⁰<https://github.com/kamangir/roofAI>

³¹`semseg`

```
semseg predict \
  [device=cpu|cuda] \
  <model-object-name> \
  <dataset-object-name> \
  <prediction-object-name>
```

3.4 Vancouver-Watching (vanwatch)

`vanwatch` ³² is a callable that,

1. terraforms the machine and the shell [2.14](#).
2. discovers the cameras in an area [3.4.1](#).
3. ingests images from the cameras discovered in an area [3.4.2](#).
4. detects the objects in the images ingested from an area and produces summary statistics [3.4.3](#).

3.4.1 vanwatch discover

Cameras are represented in different formats in different areas. `vanwatch discover area=<area>` discovers the cameras in `<area>` and stores them in `<area>.geojson` in the object `<object-name>` [2.3](#) tagged [2.1.3](#) for use by `ingest` [3.4.2](#).

```
vanwatch discover \
  [area=<area>] \
  [-|<object-name>] \
  [<args>]
```

`object-name` is tagged for retrieval by `vanwatch list discovery`.

3.4.2 vanwatch ingest

`vanwatch ingest area=<area>,count=<count> <object-name>` finds the latest set of cameras discovered [3.4.1](#) in `<area>` through tag search [2.1.3](#) and ingests `count` images into `<object-name>` and then runs `vanwatch process` [3.4.3](#) unless `~process`.

```
vanwatch ingest \
  [area=<area>,count=<-1>,model=<model-id>,<~process>,publish] \
  [-|<object-name>] \
  [<args>]
```

`object-name` is tagged for retrieval by `vanwatch list ingest`.

3.4.3 vanwatch process

`vanwatch process - <object-name>` runs object detection ³³ on the images ingested into `<object-name>` and updates `<area>.geojson`. `vanwatch process` reuses the inference in the object and completes the missing pieces.

```
vanwatch process \
  [model=<model-id>,publish] \
  [.|<object-name>] \
  [<args>]
```

+publish tags `object-name` for retrieval by `vanwatch update.QGIS`.

³²<https://github.com/kamangir/Vancouver-Watching>

³³<https://hub.ultralytics.com/models/R6nMlK6kQjSsQ76MPqQM?tab=preview>

A Concepts

A.1 Machines and Shells

A *machine* is a state machine that is connected to many other machines and shares some of its state with them for read and write. A *shell* is a stateful access mechanism to a machine that an *operator* [A.2](#) uses to run *commands* [A.3](#). Running a command in a shell can modify the state of the shell, the machine on which the shell is running, and potentially the states of all other machines.

Machines and shells may be restarted by an operator or by code. After a restart, machines and shells maintain some of their state.

Two examples of machines are a Raspberry Pi that runs Linux and is connected to AWS [34](#) and a docker container running in AWS Batch. GNU Bash [35](#) is an example of a shell.

A.2 Operators

The *operator* generates commands and runs them on different shells on different machines. The operator attempts to maximize an objective function that depends on the state of multiple machines.

A.3 Commands

A *command* is any Bash command [36](#) and can be represented in a Python string of characters [37](#). Here is an example command,

```
vanwatch ingest \
  vancouver \
  dryrun \
  . \
  --count 12
```

The above command and the one below are *identical*.

```
vanwatch discover vancouver ~upload --validate 1
```

Two commands are identical if running them on two machines in identical states yields the same states. In theory, the state of any machine depends on the state of any other machine, and it is almost impossible to run two commands in identical states, including the time of execution. Therefore, when we refer to two identical commands, we either use a derivation-based proof of identity or consider a validation in a limited “relevant” subset of the state representation.

For any shell on any machine, at known states, there is a mapping between the set of all commands and $\{True, False\}$ that we address as “whether the command is found”. In Bash, for example, the following message is printed when a command “is not found”.

```
-bash: void: command not found
```

Note that writing to the standard streams *stdin* and *stdout* are examples of state changes in the shell and the machine.

Terraforming is the process of running commands that modify the state of the shell and the machine in ways that make additional commands found. Terraforming is also intended to modify the state change caused by a set of commands favourable to the interests of an operator. For convenience, we address a command “that is found” as a *valid* command and *invalid* otherwise. Terraforming may also ensure the states of the machine and the shell after a restart. Terraforming generally includes a modification of *bashrc*, *bash_profile*, *screenrc*, and *desktop* files.

Commands know the machine and the shell they are running in and adjust their operation accordingly. For example, a script that submits jobs inside a docker container may download the artifacts generated through previous submissions on a user-facing machine, such as a Macbook.

³⁴<https://aws.amazon.com/>

³⁵<https://www.gnu.org/software/bash/>, *Bash*, *bash*.

³⁶<https://www.gnu.org/software/bash/manual/bash.html#Shell-Syntax>

³⁷<https://docs.python.org/3/library/string.html>

The first word ³⁸ in a command is generally the callable A.3.1. The rest of the command is expected to follow the conventions of the callable. In this paper we propose guidelines that we later demonstrate lead to useful expansions 2.

A.3.1 Callables

A *callable* is a valid command with no space and control operators ³⁹. The list of callables depends on the machine’s state and is generally extended through terraforming. Some of the well-known callables are *git* ⁴⁰, *docker* ⁴¹, *pushd* ⁴², *nano* ⁴³.

Theorem 1 For any callable $\langle \text{callable} \rangle$, and any string $\langle \text{suffix} \rangle$, “ $\langle \text{callable} \rangle \langle \text{suffix} \rangle$ ” is a valid command.

Theorem 2 For any valid command $\langle \text{command} \rangle$, and any string $\langle \text{suffix} \rangle$, “ $\langle \text{command} \rangle \langle \text{suffix} \rangle$ ” is a valid command.

The `core 2.9` is a callable. Many plugins 2.10 define their callable.

A.3.2 Command Templates

Commands can be similar when considered as strings of characters. Here is a command that is similar to the above,

```
vanwatch ingest toronto upload . --count 3
```

A *command template* is a representation that yields similar commands, given the following rules. First, $\langle \text{description} \rangle$ can be replaced with any string of characters that can be described as “description”. See *objects 2.3*, *options 2.2*, and *arguments ??*, for the next rules. Here is a command template for the two above,

```
vanwatch ingest \
  <area> \
  [dryrun,~upload] \
  [<object-name>] \
  [--count <-1>]
```

A.4 Computational Model

A group of operators create and modify a set of scripts 2.11 maintained in a code repository, such as *git 2.15*. The `core 2.9` and the plugins 2.10 are also maintained in one or more repositories. Each operator can access a set of machines and create shells on them. Each operator can also access a set of repositories and clone them on the machines where they create shells and receive updates. The operator can modify any of the repositories that they have access to following a collective peer-reviewed *pull-request* ⁴⁴ process.

The operators act asynchronously while communicating with each other. Multiple operators may simultaneously use the same machine, and the same operator may simultaneously use multiple machines. Only one operator uses a shell at one time.

Some machines are exogenous to this model, yet the operators can access their states in read or write mode. Cloud storage and compute resources are examples of these machines.

³⁸https://www.gnu.org/software/bash/manual/html_node/Shell-Syntax.html

³⁹<https://www.gnu.org/software/bash/manual/bash.html#index-control-operator>

⁴⁰<https://git-scm.com/docs/git>

⁴¹<https://docs.docker.com/engine/reference/commandline/cli/>

⁴²<https://www.gnu.org/software/bash/manual/bash.html#Directory-Stack-Builtins>

⁴³<https://www.nano-editor.org/>

⁴⁴<https://docs.github.com/en/pull-requests/collaborating-with-pull-requests/proposing-changes-to-your-work-with-pull-requests/about-pull-requests>

B Background and Context

Almost five years ago, on Thursday, November 8, 2018, I acquired a Raspberry Pi ⁴⁵ on Amazon. Since then, my personal and professional lives have focused on Linux. Professionally, I do AI and, more recently, geospatial AI. In my personal life, I mix AI, cloud, and mathematics into minimal forms that seek survival ⁴⁶. Over the years, I have built a set of mechanisms for building AI systems that I will document in this paper. Therefore, this is an attempt to produce formal mathematical definitions for the AI mechanisms that I will collectively refer to as *giza*. I seek to understand these mechanisms through this effort better to use them more optimally and along new dimensions.

This paper discusses concepts at the intersection of mathematics, software science, and computer science, and lacks scientific rigour in many places. I intend to push the practical development of this theory to fruition and hope to receive guidance along the way from experts in the field and solidify the theoretical foundations.

⁴⁵<https://www.raspberrypi.org/>

⁴⁶<https://github.com/kamangir>